

Multithreaded Sorting

Topic - Implementation and analysis of multithreaded sorting where multiple threads are used to sort segments of an array and these segments are merged (with and without parallelism)

Submitted By: Kushagra Indurkha
CS19B1017

Contents

Implementation Details	2
Multithreaded sorting of Segments of Array	2
Method-1: Sequential merging in main thread	5
Method-2: Multithreaded Merging	6
Time Graph Analysis	9
Graph-1	10
Graph-2	11
Exploring further	12
Conclusion	13

Implementation Details

As stated in the [problem statement](#) to study the efficiency of using multithreading in performing sorting of arrays i have implemented 2 multithreaded solutions which take (n,p) as input where 2^n is the length of array and 2^p is the number of threads.

Multithreaded sorting of Segments of Array

- The main thread(master) creates 2^p slave threads using `pthread_create` each sorting a segment of length $2^{(n-p)}$ of the array

```
for (int i = 0; i < no_of_threads; ++i)

pthread_create(&thread_ids[i],NULL,sort_seg, (void*)&shared_data[i]);
```

Thread Creation

The following attributes are passed in `pthread_create` call :

- `&thread_ids[i]` - `thread_ids` is an array of type `pthread_t` and length 2^p to store the IDs of all the threads created.
- `NULL` - to use the default attributes
- `Sort_seg` - function calls the sort function which is implementation of standard merge sort algorithm called on the segment of array in range(`start`,`start+seg_size`) and exits with `NULL` value when the control returns after sorting the segment.

```
void* sort_seg(void* seg_start)
{
    share* data=(share*)seg_start;
    int start=data->start;
    long int* arr=data->arr;
    int seg_size=data->seg;
```

```

•   int len=data->len;
•
•   sort(arr,start,start+seg_size-1,len);
•   pthread_exit(NULL);
•
• }

```

This function in turn calls the `sort` function which is a recursive merge sort implementation which takes array and left and right indices of the sub array to be sorted , the `merge` function merges 2 sorted subarrays boundary indices of which are passed as arguments to it:

```

void merge(int a,int b,int c,int d,long int* arr)
{
    int temp[d-a+1];
    int i=a,j=c,k=0;
    while(i<=b && j<=d)
    {
        if (arr[i]<arr[j])
            temp[k++]=arr[i++];
        else
            temp[k++]=arr[j++];
    }
    while (i<=b)
        temp[k++]=arr[i++];
    while (j<=d)
        temp[k++]=arr[j++];

    for (int i = a,j=0; i <=d; i++,j++)
        arr[i]=temp[j];
}

```

```

void sort(long int*arr,int l,int r)
{
    if(l>=r)
        return;

```

```

    int m = (l+r-1)/2;
    sort(arr,l,m);
    sort(arr,m+1,r,len);
    merge(l,m,m+1,r,arr,len);
}

```

- `&shared_data[i]` - is the argument passed to the `sort_seg` function which is a pointer to the *i*'th element of `shared_sata` array of a struct data typecasted to a void pointer, the composition of the struct is :

```

long int* arr -> pointer to the array of random long ints
int start -> start index of the segment
int seg      -> length of the segment (2^(n-p))

```

- After all the threads are created and with each thread sorting a segment of the array , the main(master) thread wait for all the slave to terminate this is achieved by using `pthread_join()`

```

for (int i = 0; i < no_of_threads; i++)
    pthread_join(thread_ids[i], NULL);

```

Here we wait for every thread ID in `thread_ids` array to finish

Now `arr` has its segments (length= $2^{(n-p)}$) sorted, now we merge those segments in two different ways:

Method-1: Sequential merging in main thread

We will use main thread to merge our segments in a sequential manner for example if `no_of_threads=8` and there are 8 segments numbered 1 to 8 then the merging will proceed as follows:

Phase-1	(1)(2)
Phase-2	(1,2)(3)
Phase-3	(1,2,3)(4)
Phase-4	(1,2,3,4)(5)
Phase-5	(1,2,3,4,5)(6)
Phase-6	(1,2,3,4,5,6)(7)
Phase-7	(1,2,3,4,5,6,7,8)

(Segments in '()' are already merged)

This is my implementation of merging two sorted subarrays of length `p,q` which has both time and space complexity of $O(p+q)$

```
void merge(int a,int b,int c,int d,long int* arr,int len)
{
    int temp[d-a+1];
    int i=a,j=c,k=0;
    while(i<=b && j<=d)
    {
        if (arr[i]<arr[j])
            temp[k++]=arr[i++];
        else
            temp[k++]=arr[j++];
    }
    while (i<=b)
        temp[k++]=arr[i++];
    while (j<=d)
        temp[k++]=arr[j++];
}
```

```

for (int i = a, j=0; i <=d; i++,j++)
    arr[i]=temp[j];
}

```

To merge the segments in the manner mentioned above i have called the merge function in the following manner:

```

for (int i = 1; i < no_of_threads; i++)

merge(0, (i*seg_size)-1, i*seg_size, (i*seg_size)-1+seg_size, &arr[0], len);

```

After $(2^p)-1$ iterations all the sorted segments are merged and the whole array is sorted

Method-2: Multithreaded Merging

In method-2 i have taken a multithreaded approach in merging the sorted segments of the array in such way that the main thread creates $\text{no_of_threads}/2$ slave threads where each thread merge two segments and exits then in the next iteration $\text{no_of_threads}/4$ slave threads are created to merge the merge the merged segments and so on until only one segment remains.

For example if $\text{no_of_threads}=8$ and there are 8 segments numbered 1 to 8 and corresponding threads numbered (0 to 7) then the merging will proceed as follows:

Phase-1: 1,2 3,4 5,6 7,8
Phase-2: 1,2,3,4 5,6,7,8
Phase-3: 1,2,3,4,5,6,7,8

(Different colors indicate different threads)
Thread-0, **Thread-2**, **Thread-4**, **Thread-6**

To create the slave threads to merge the segments in the manner mentioned above i have used `pthread_create` called the in the following manner:

```
1.     merge_thread_data data[no_of_threads];
2.     for (int i = 1; 2*i<=no_of_threads; i*=2)
3.     {
4.         int j=0;
5.         for (int thread_idx = 0; thread_idx < no_of_threads ;
6.             thread_idx++)
7.         {
8.             if (thread_idx%(2*i)==0)
9.             {
10.                 data[thread_idx].arr=&arr[0];
11.                 data[thread_idx].arr_len=len;
12.                 data[thread_idx].phase=i;
13.                 data[thread_idx].seg_size=seg_size;
14.                 data[thread_idx].seg_idx=thread_idx;
15.
16.                 pthread_create(&thread_ids[j],NULL,merge_thread, (void*)&data[thread_idx]);
17.                 j++;
18.             }
19.         }
20.         for (int k = 0; k < j; k++)
21.             pthread_join(thread_ids[k], NULL);
22.     }
```

Here the control variable of outer loop (line 2 to 20) is multiplied by 2 every time it indicates the no of consecutive array segments which are currently merged,in case of the given example this runs for 3 times(3 phases)

Then in the inner loop we check if the segment corresponding to i needs to be processed in this iteration by taking its modulus with 2*i.

(To prevent the time lost in array creation i have reused the same thread array used during sorting).

Variable j is incremented each time a slave thread is created

The following attributes are passed in `pthread_create` call :

- `&thread_ids[i]` - `thread_ids` is an array of type `pthread_t` and length 2^p to store the IDs of all the threads created.
- `NULL` - to use the default attributes
- `Merge_thread` - Calls the merge function (same as that of method-1) with corresponding subarray's starting and ending indices And exits after they are merged

```
void* merge_thread(void* shared_data)
{
    merge_thread_data* data=(merge_thread_data*) shared_data;
    long int* arr=data->arr;
    int len=data->arr_len;
    int phase=data->phase;
    int seg_size=data->seg_size;
    int seg_idx=data->seg_idx;

    int a=seg_idx*seg_size;
    int b=a+(phase*seg_size)-1;
    int c=b+1;
    int d=c+(phase*seg_size)-1;

    merge(a,b,c,d,arr,len);

    pthread_exit(NULL);
}
```

- `&data[thread_idx]` - is the argument passed to the `merge_thread` function which is a pointer to the element of data array of `merge_data` struct typecasted to a void pointer, the composition of the struct is :


```

struct merge_data
{
    long int* arr    -> pointer to the array of random long ints
    int seg_idx      -> start index of the segment
    int seg_size     -> length of the segment (2^(n-p))
    int arr_len      -> length of the array (2^n)
    int phase        -> i ie the phase of merging
}

```

After the threads are created in the inner loop before exiting the outer loop all the created threads are joined (line-18,19).

In this way all the sorted segments are merged.

Time Graph Analysis

To analyse the efficiency of the above described methods , running three solutions (Solution-1 Method-1, Solution-2 Method-2, Solution-3 is the plain usage of merge sort algorithm) with varying n,p gave the following results:

```

stark@Jarvis:/mnt/d/WSL/os_assignment/submission$ ./visualize.sh
✓ Compiled
✓ Inputs Generated

```

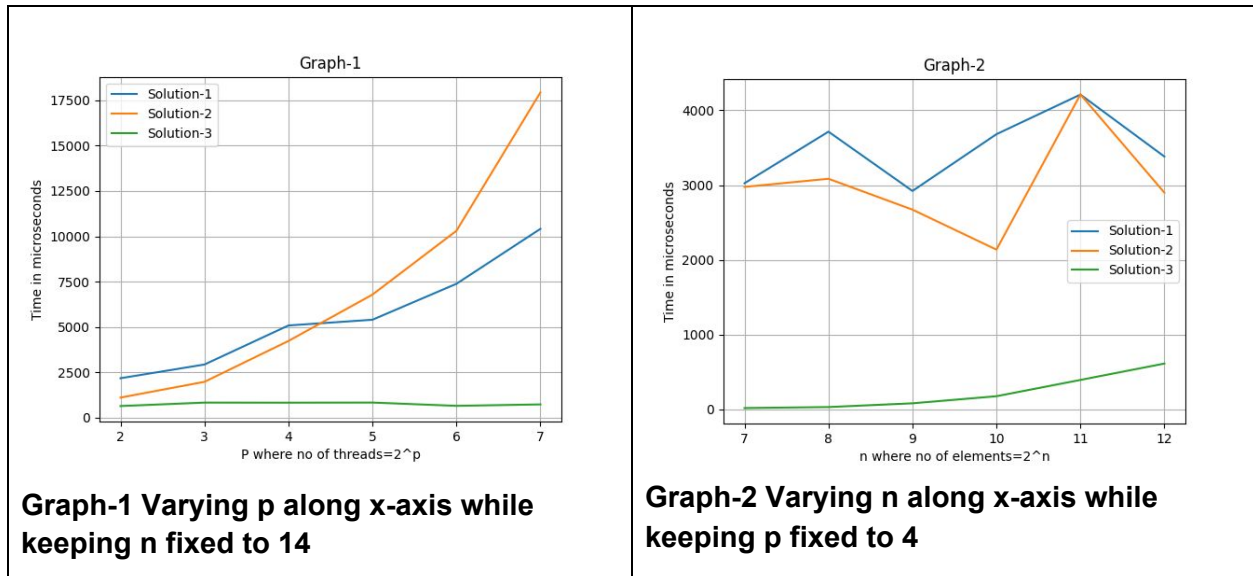
			Graph-1			
(n p)	SOL-1		SOL-2		SOL-3	
(12 2)	2179	SORTED	1112	SORTED	646	SORTED
(12 3)	2935	SORTED	1981	SORTED	836	SORTED
(12 4)	5087	SORTED	4233	SORTED	832	SORTED
(12 5)	5403	SORTED	6788	SORTED	838	SORTED
(12 6)	7384	SORTED	10305	SORTED	657	SORTED
(12 7)	10412	SORTED	17932	SORTED	736	SORTED

			Graph-2			
(n p)	SOL-1		SOL-2		SOL-3	
(7 4)	3025	SORTED	2976	SORTED	18	SORTED
(8 4)	3715	SORTED	3085	SORTED	31	SORTED
(9 4)	2922	SORTED	2672	SORTED	81	SORTED
(10 4)	3681	SORTED	2137	SORTED	176	SORTED
(11 4)	4211	SORTED	4213	SORTED	394	SORTED
(12 4)	3382	SORTED	2899	SORTED	613	SORTED

```

✓ Graphs Plotted
✓ Cleanup
Script Exited

```



Graph-1

Observations:

1. Time taken by Both solution-1 and solution-2 increase as we increase the number of threads
2. Solution-2 performs better than Solution-1 for $p=2,3,4$.
3. Solution-3 performs better than both solutions-1 and 2.

Explanations:

1. As we increase the number of threads the ,the thread overhead (caused by thread creation, deletion, scheduling, management and context switches) is more than the time saved by the use of multithreading in sorting the segments
2. Initially when no of threads is less the overhead caused by thread creation , deletion and context switches (in merging) is compensated by the speedup provided by the multithreading, As we increase the number of threads the time taken by the thread overhead is more than the time saved by the use of multithreading in merging the segments.
3. In a relatively small array the thread overhead is much more than the speedup provided by multiple threads(both in sorting and sorting+merging) hence plain merge sort performs better than both.

(Note- observation-2,3 are corresponding to the graph above , its not true for every execution as many other factors also control the running time)

Graph-2

Observations:

1. Solution-2 in general performs better than Solution-1.
2. Solution-3 performs better than both solutions-1 and 2.
3. The trends in solution-1 & 2 are erratic

Explanations:

1. As we increase the number of elements the work given to each thread is increases (large segments merging) hence the thread overhead is compensated and justified also the no of user threads fixed here is not very large than the no of kernel threads available hence parallelism is used aptly without unnecessary thread creation, management, and scheduling overhead while in solution 1 multithreading is used only in sorting and no parallelism is used for merging hence solution 2 performs better than solution 1.
2. In the range of n from 7 to 12 no of elements are from 128 to 4096 which is relatively a small size here the thread management overhead cannot be compensated by the speedup provided by multithreading (both in sorting and sorting+merging) hence plain merge sort performs better than both.
3. The trends in solution-1 and solution-2 are erratic as the time taken by both the solutions is dependent on various other factors such as the availability of cpu resources, execution time of threads and other processes in the queue.

Exploring further

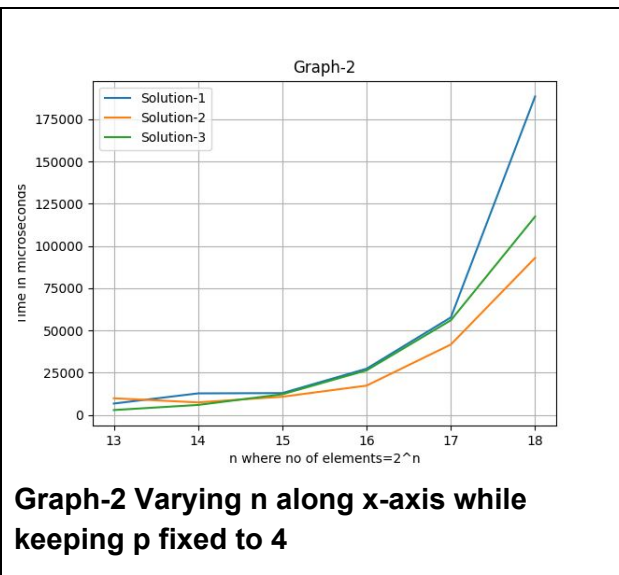
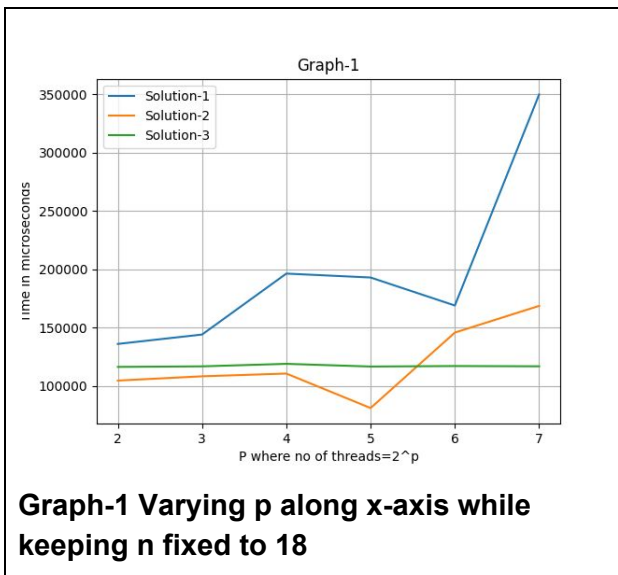
Checking on larger values of n:

```
stark@Jarvis:/mnt/d/WSL/os_assignment/submission$ ./visualize.sh
✓ Compiled
✓ Inputs Generated

Graph-1
(n p)  SOL-1      SOL-2      SOL-3
(18 2) 136042    104628    116357    SORTED
(18 3) 144101    108285    116822    SORTED
(18 4) 196365    110637    118975    SORTED
(18 5) 192995    81076     116669    SORTED
(18 6) 168995    145792    117090    SORTED
(18 7) 349977    168669    116861    SORTED

Graph-2
(n p)  SOL-1      SOL-2      SOL-3
(13 4) 6805       9911       2964      SORTED
(14 4) 12822      7487       5979      SORTED
(15 4) 12991      10843      12294     SORTED
(16 4) 27424      17411      26455     SORTED
(17 4) 57769      41735      56063     SORTED
(18 4) 188471     92993      117437    SORTED

✓ Graphs Plotted
✓ Cleanup
Script Exited
```



- In graph-1 we observe that solution-2 is performing better than plain mergesort and solution-1 when p is varying from 2 to 5.
- In graph-2 we observe that solution-2 is performing better than both solution-1 & 3 for n ranging from 15-18.

Conclusion

After running all 3 solutions on different sets of values various times i conclude the following:

- For smaller arrays (approximately of size ≤ 20000) using simple merge sort without parallelism is efficient.
- For large arrays (order of 10^5) using multithreading for sorting as well as merging (Solution-2) is most efficient.
- For small arrays, if using solution-1 or 2 the value of p should be kept minimum - {1,2} so that the number of threads are {2,4}, this can sometimes give even better performance than plain mergesort.
- For large arrays value of p can be kept around {4,5} for optimal performance.